

Messaging

in distributed systems

micro-23

Aliaksei Bialiauski

Designed in L^AT_EX

All visual and text materials presented in this slidedeck are either originally made by the author or taken from public Internet sources, such as website. Copyright belongs to their respected authors.



Request-Response vs Producer-Consumer

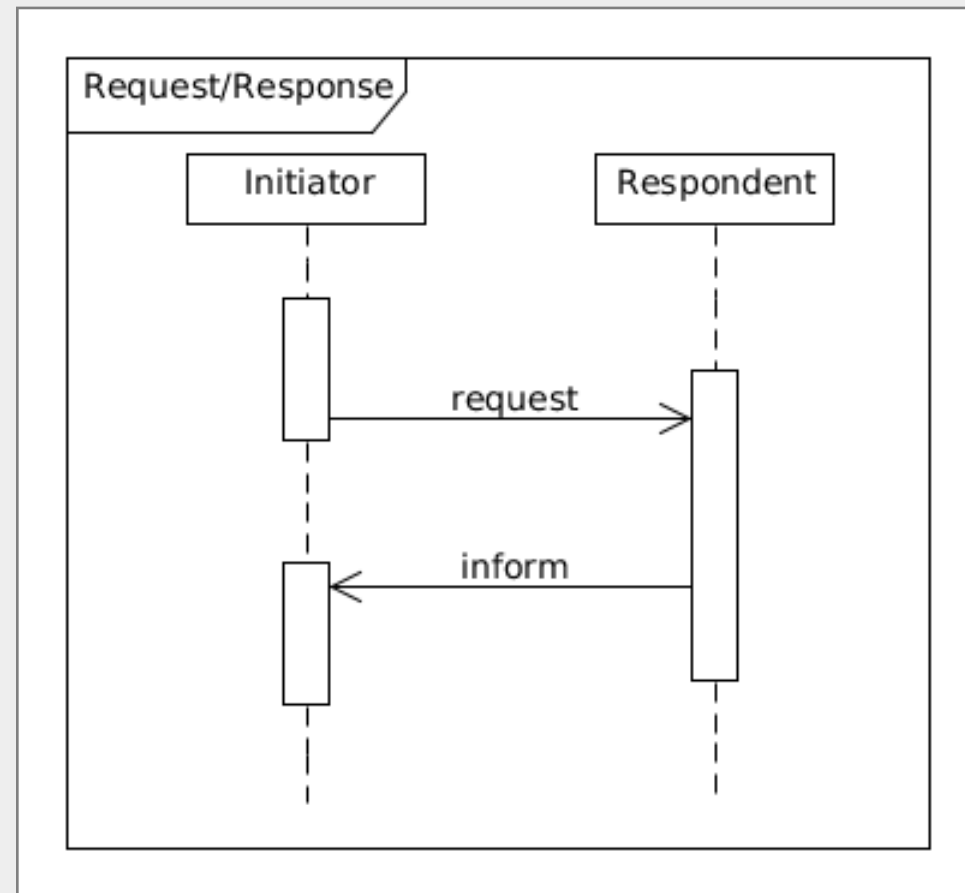
Messaging Patterns

Internals of messaging

Chapter #1:

Request-Response vs Producer-Consumer

Request-Response



Availability

What should the client do when the server is not available?

Performance

What should the client do when the server gets slower?

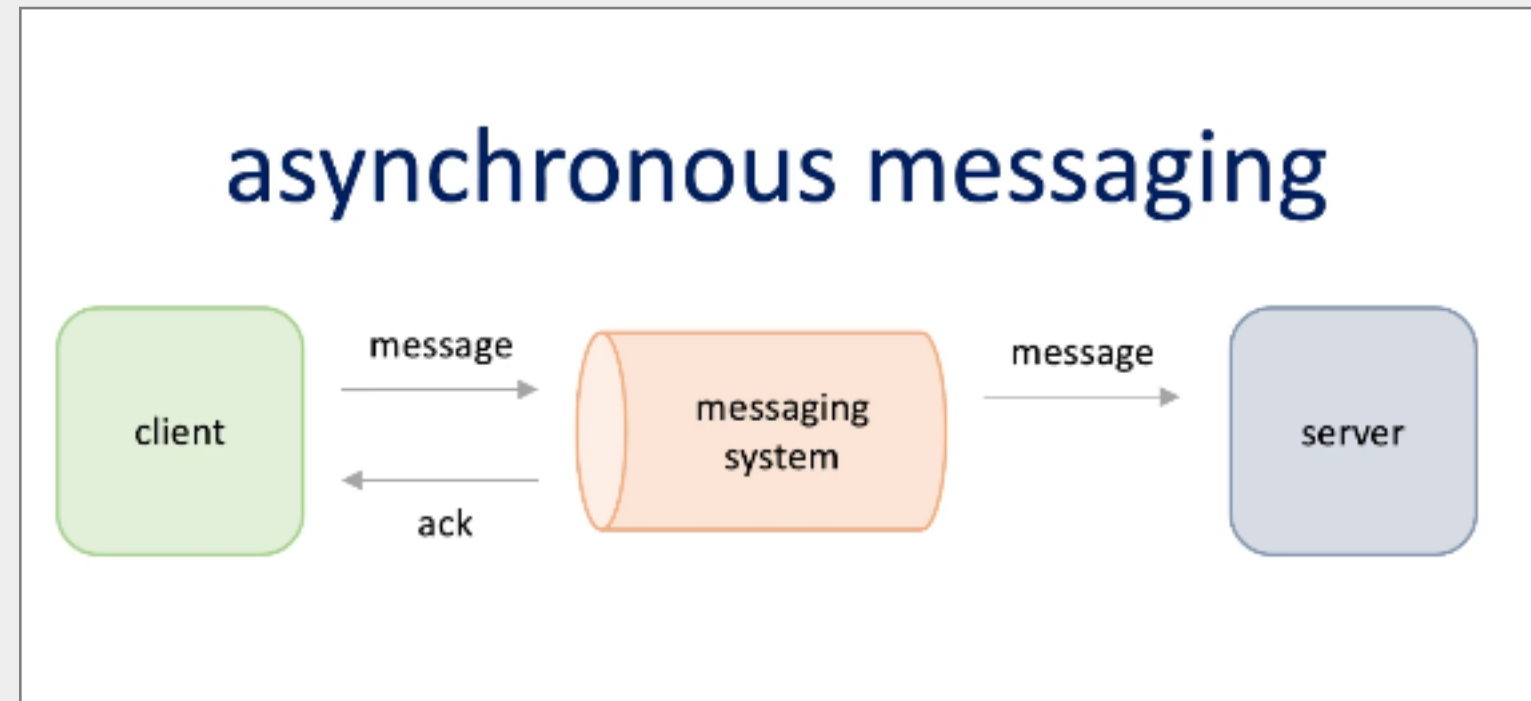
Backpressure

What should the server do when there is a sudden traffic spike?

Retry

What should the client do with failed requests?

Asynchronous messaging



[Availability Performance Backpressure Retry [Async](#)]

Availability -> Keep messages and send them later

Performance -> Add more servers to parallelize processing

Backpressure -> Keep draining messages at its own pace

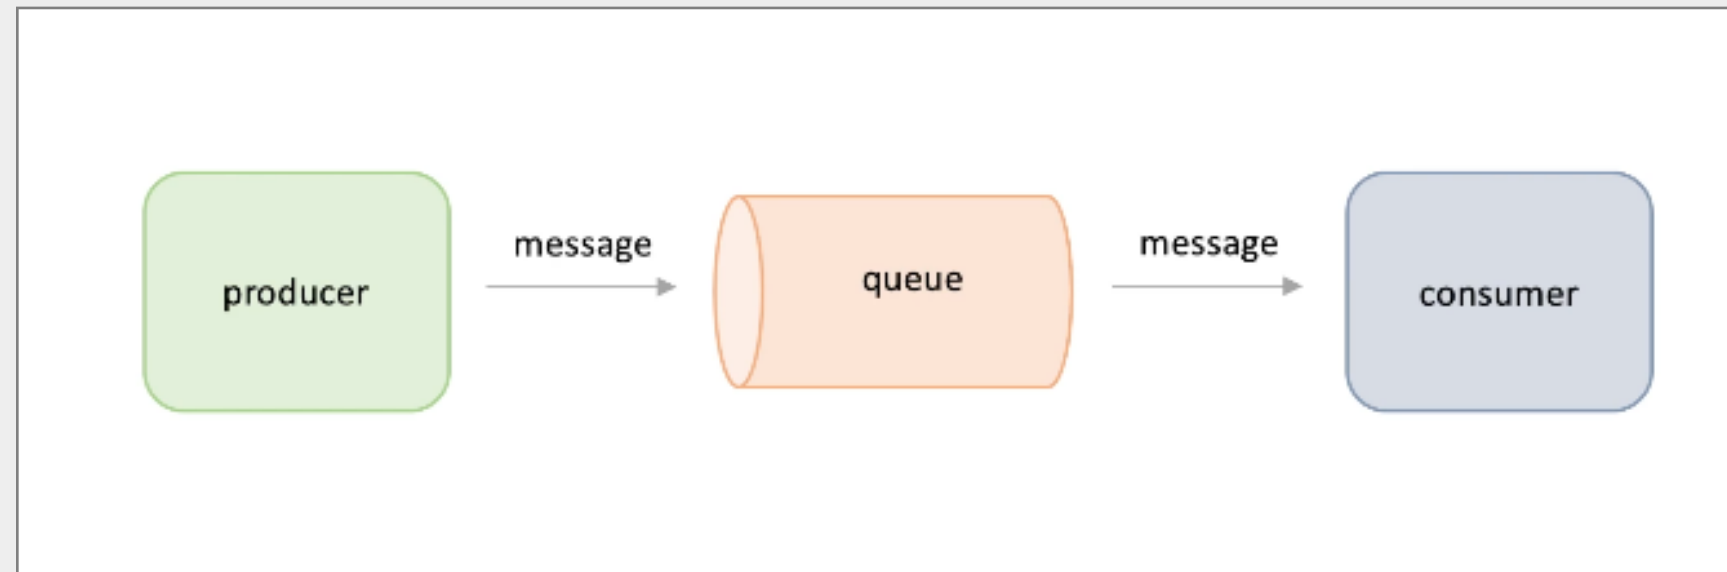
Retry -> Re-send failed messages

Also, asynchronous messaging helps to **decouple** client and server.

Chapter #2:

Messaging Patterns

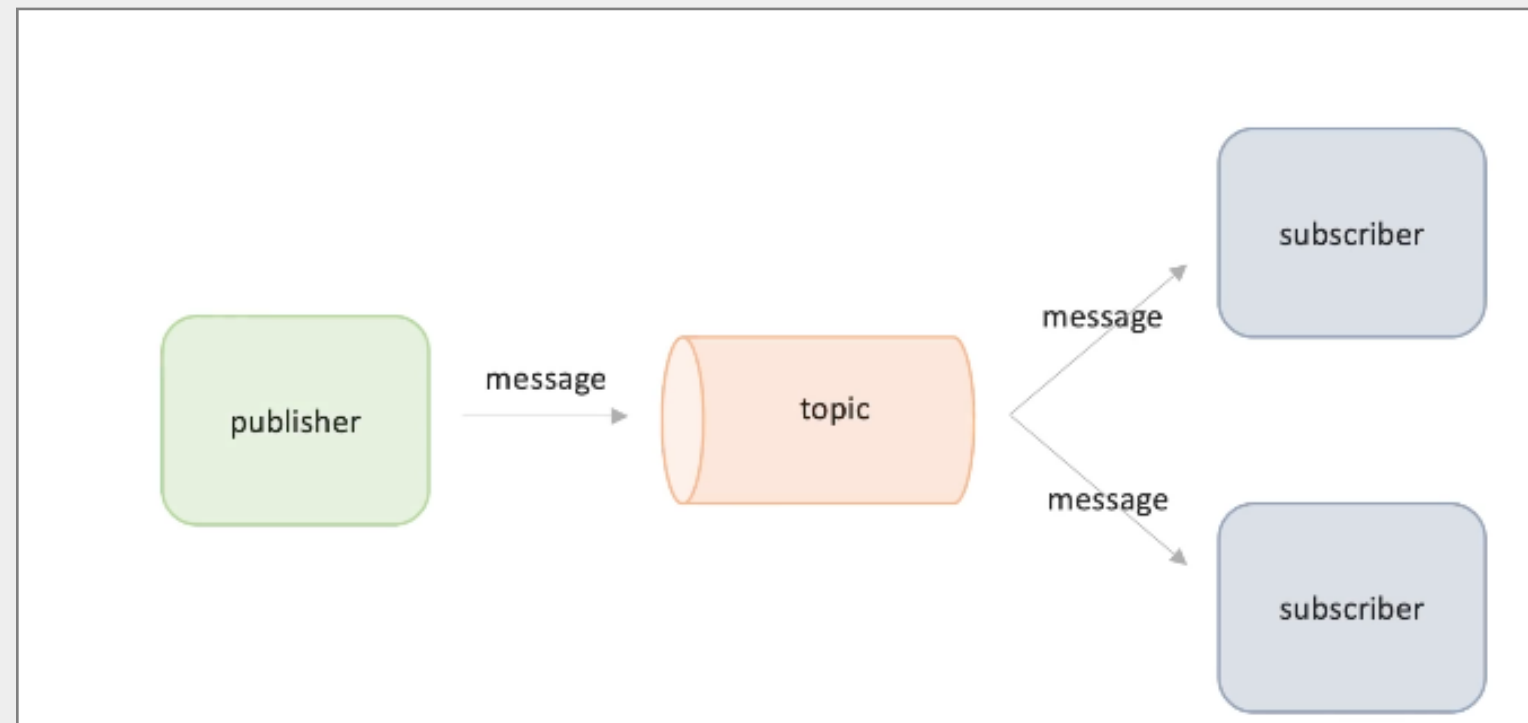
Message Queuing



<https://aws.amazon.com/sqs>

<https://www.rabbitmq.com>

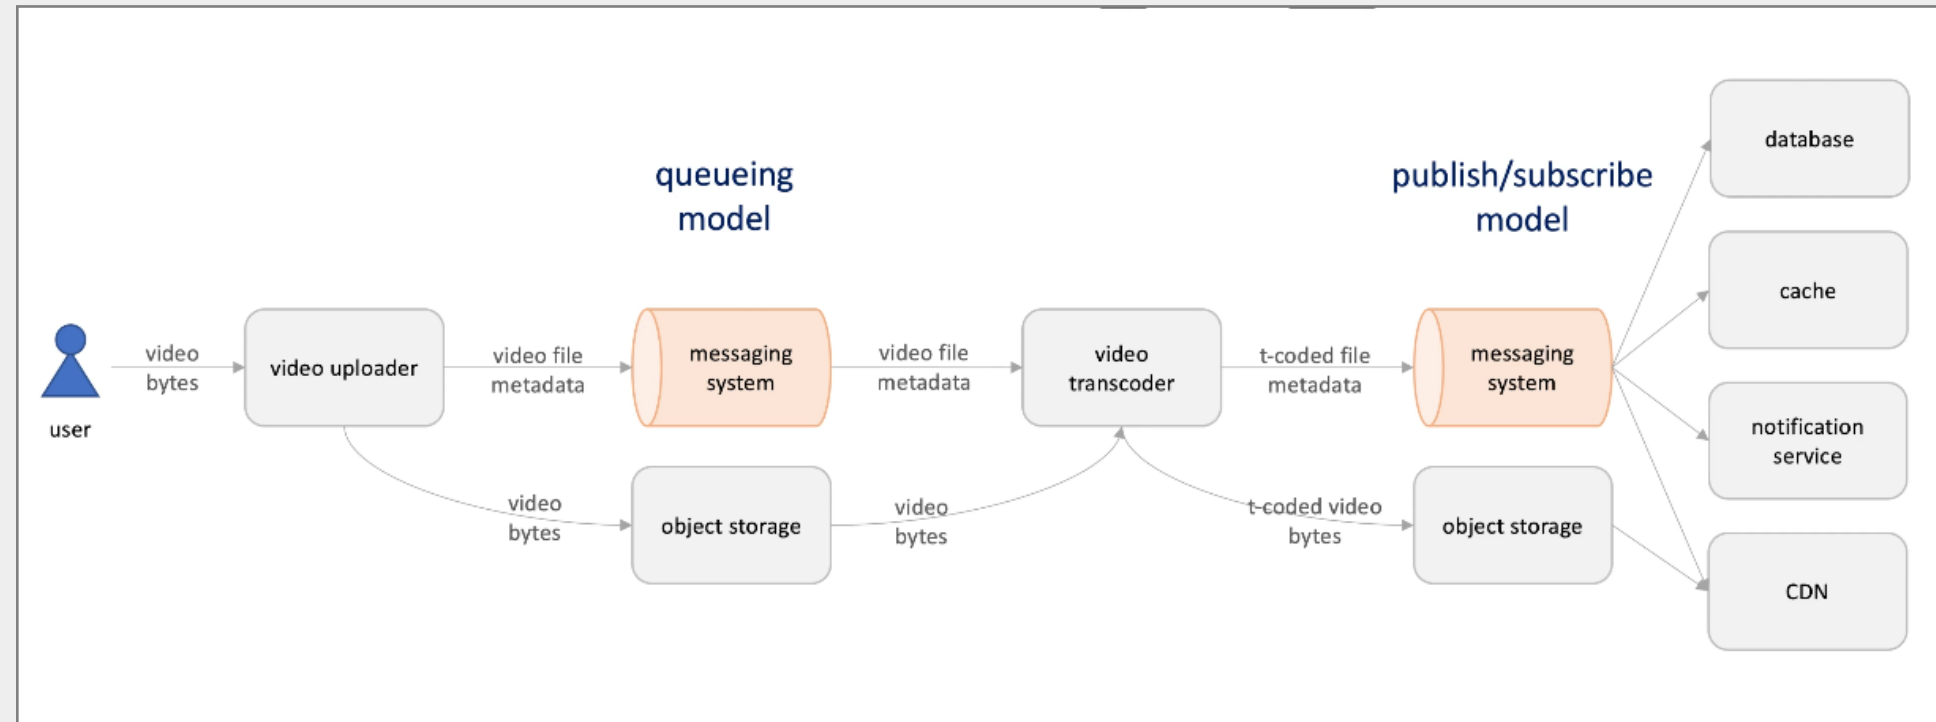
Publish-Subscribe



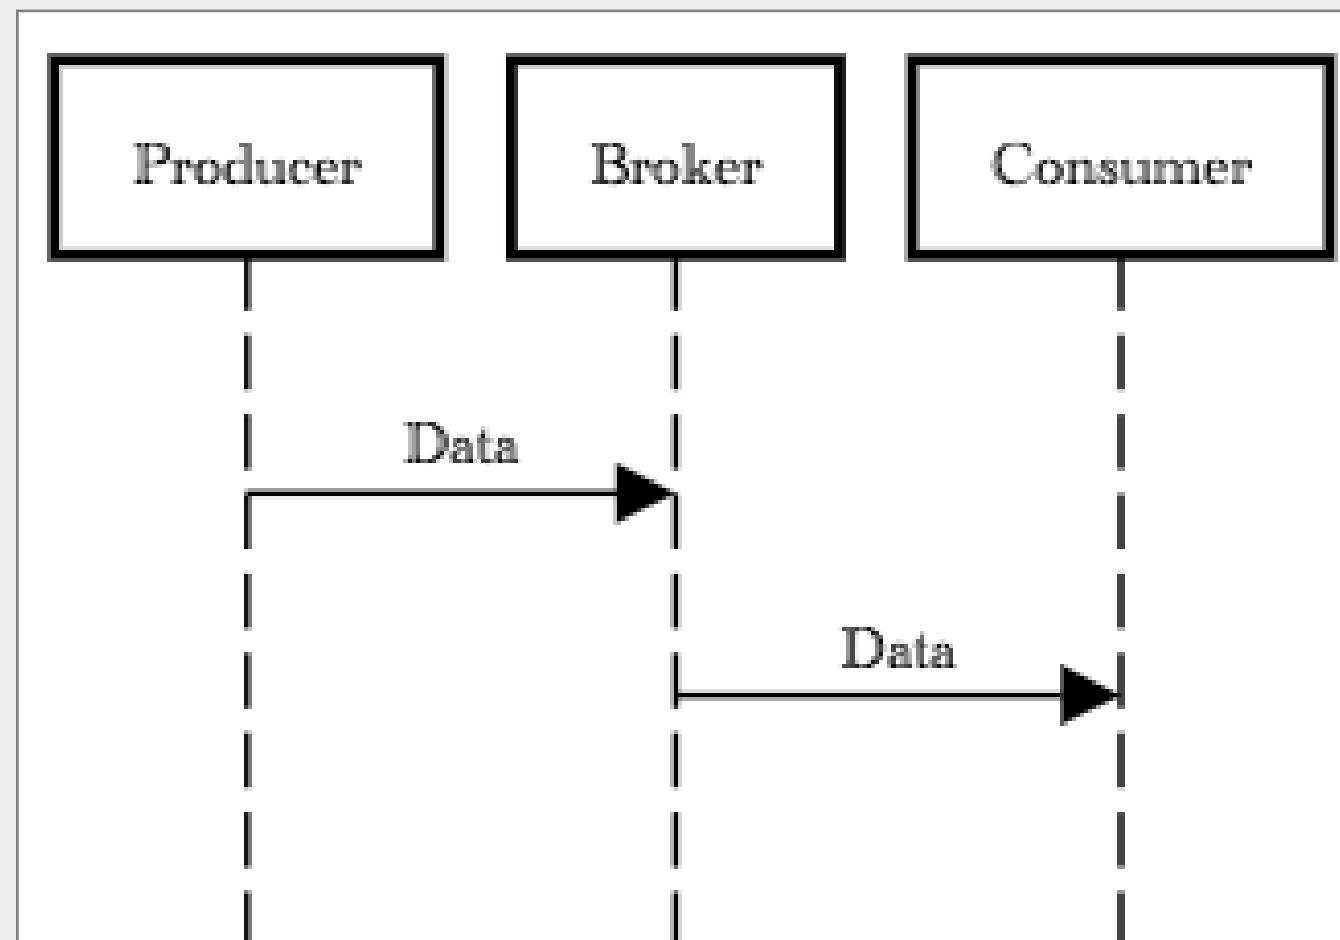
<https://kafka.apache.org>

<https://aws.amazon.com/kinesis>

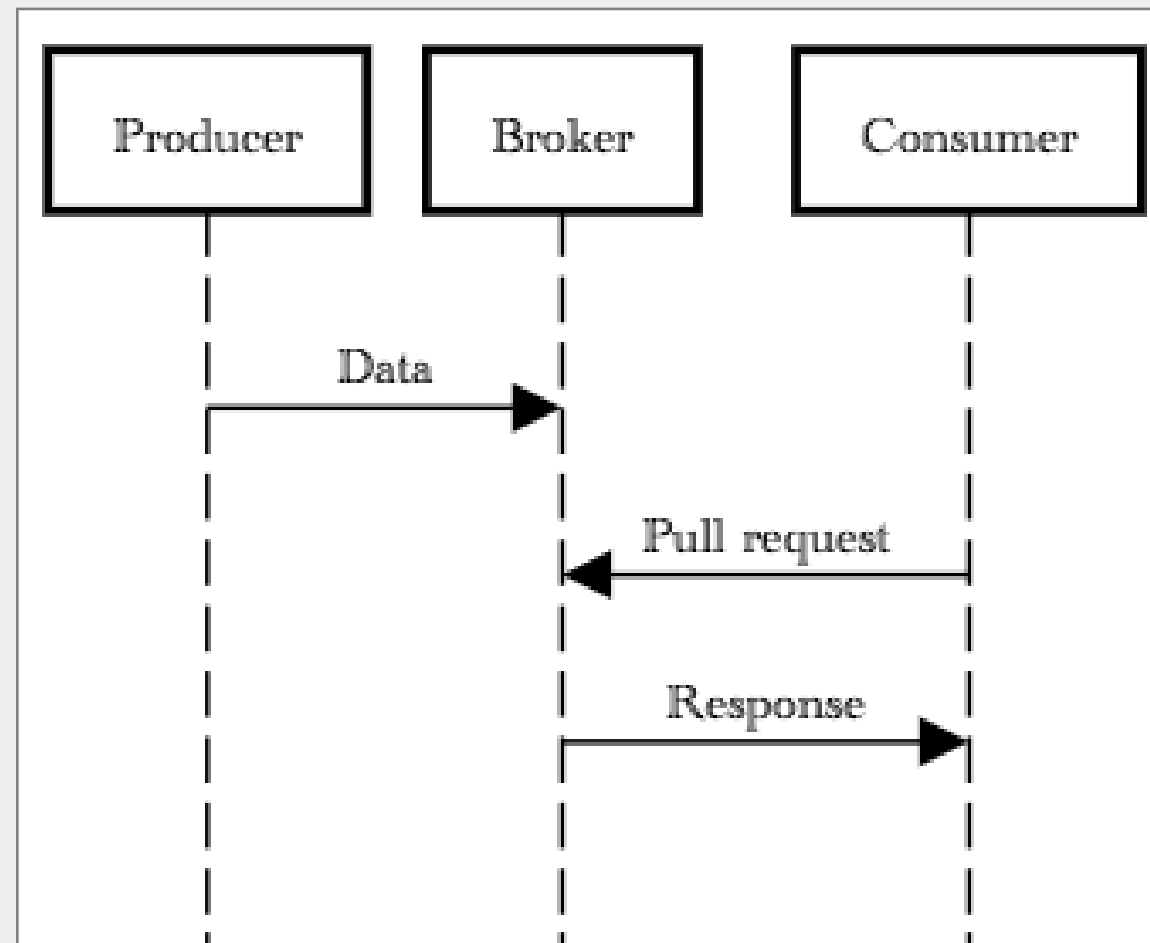
[MQ Pub-Sub P&P Competing Claim Retry Partitioning]



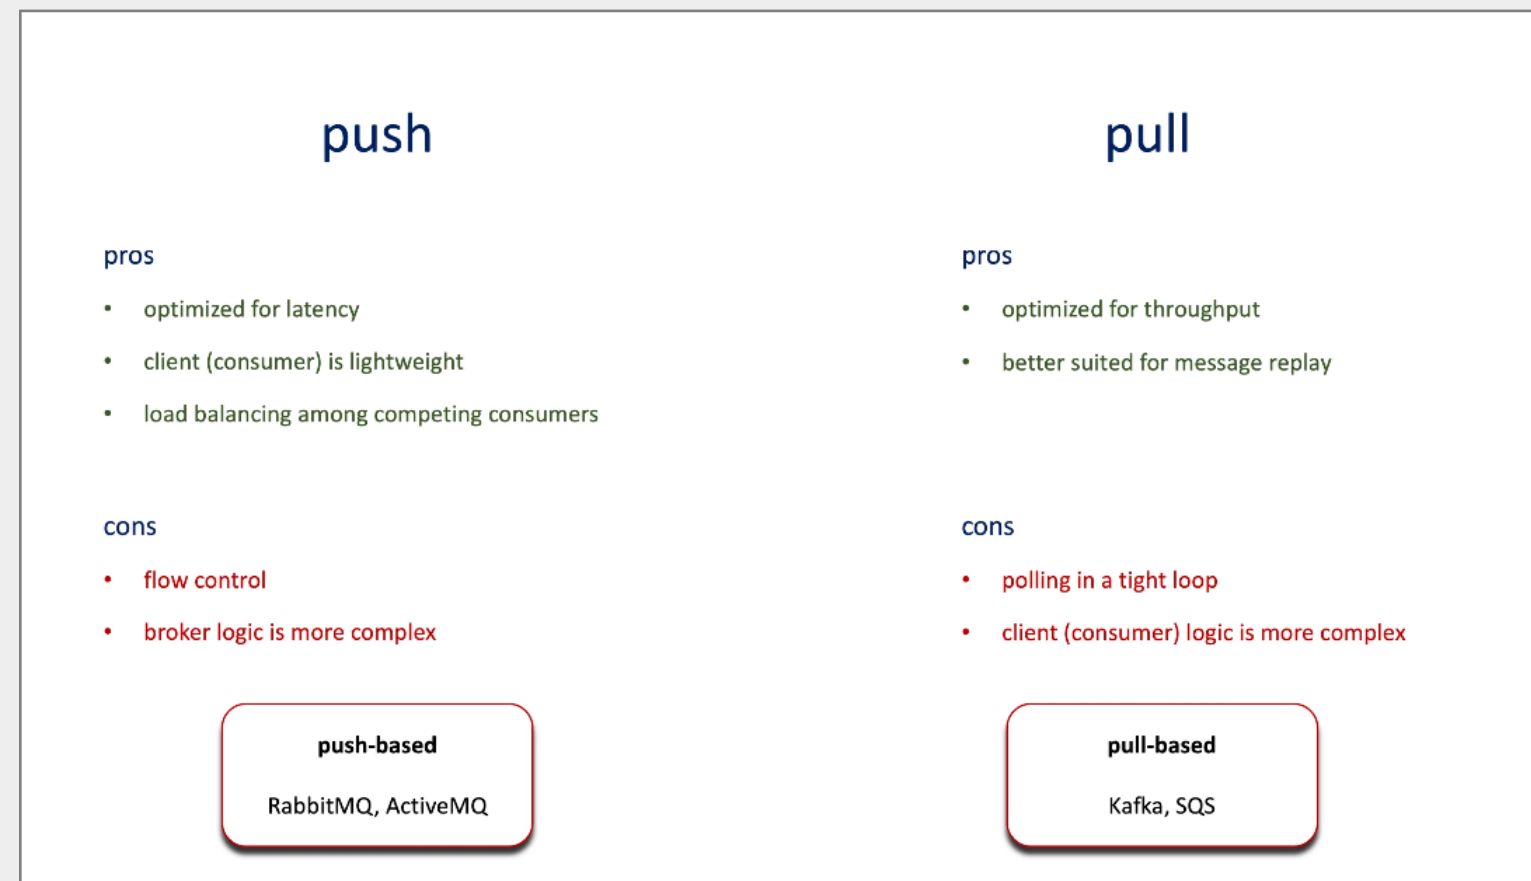
Push model



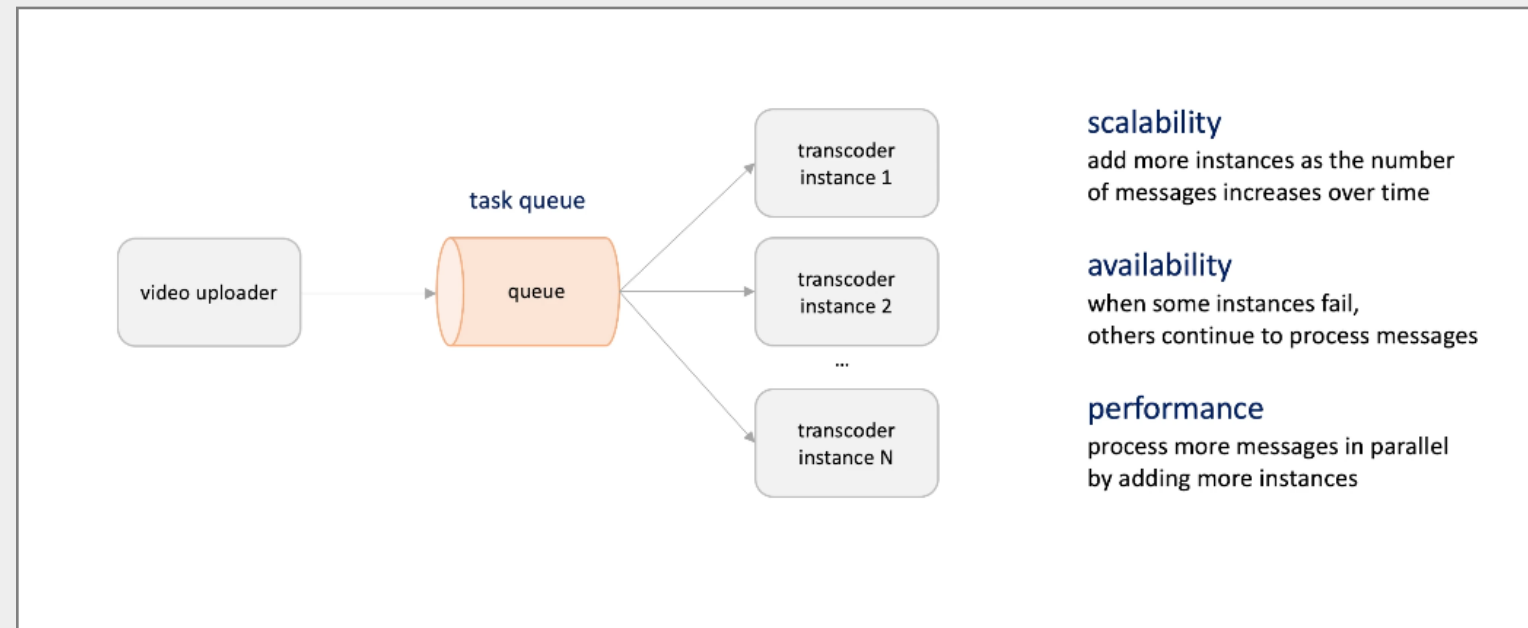
Pull model



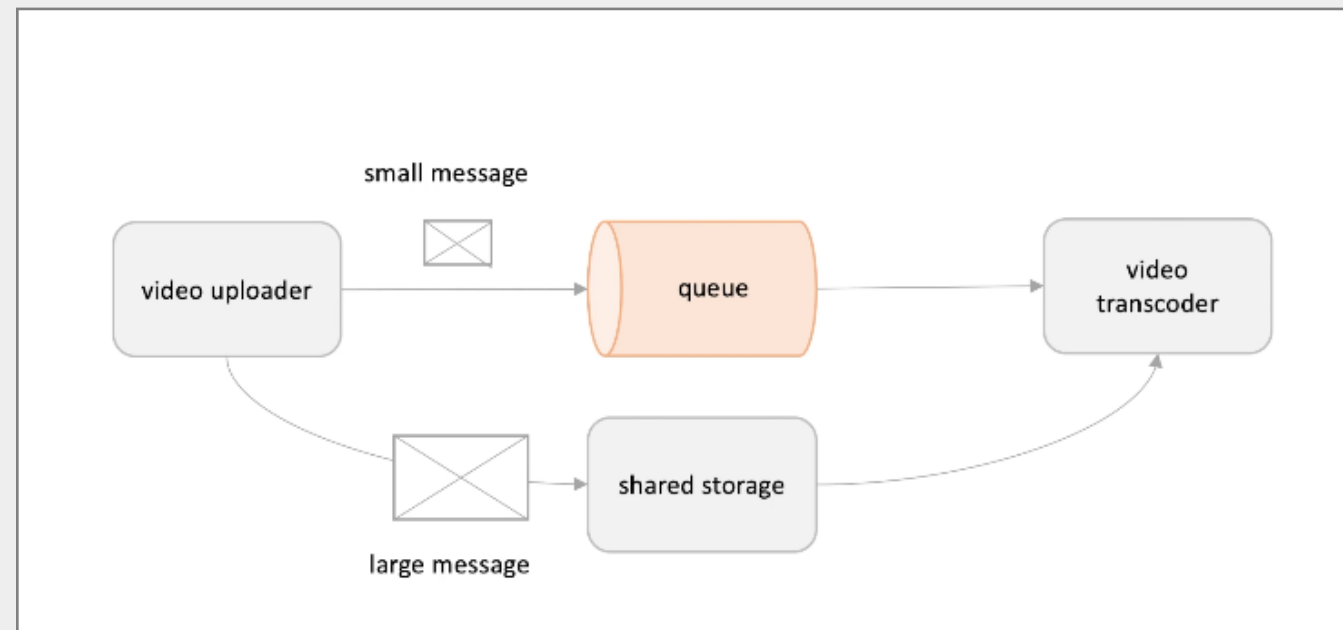
Push vs Pull



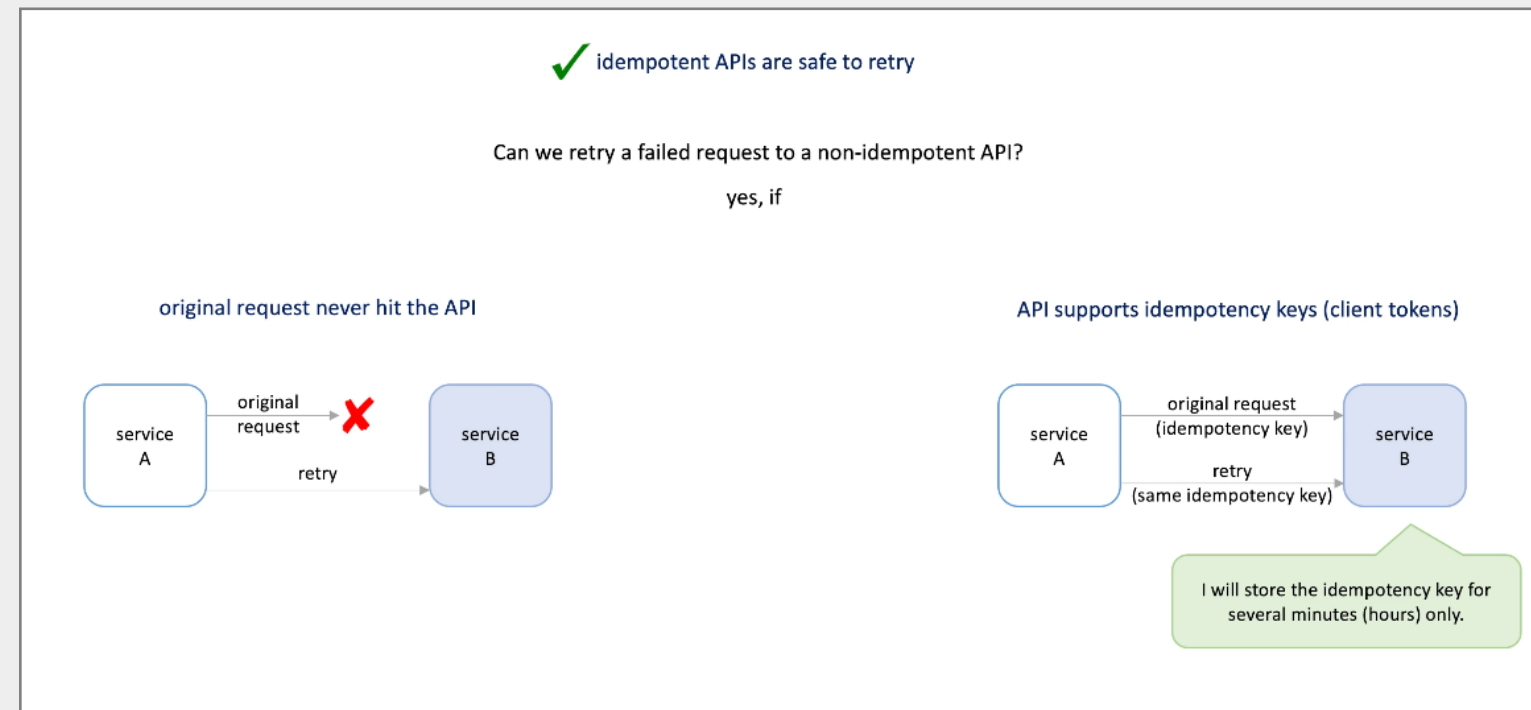
Competing Consumers



Claim Check

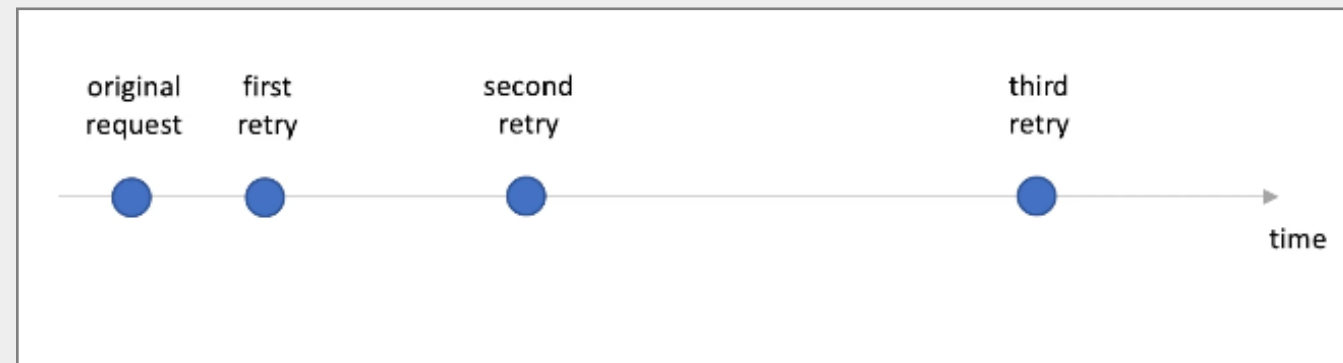


When to retry?

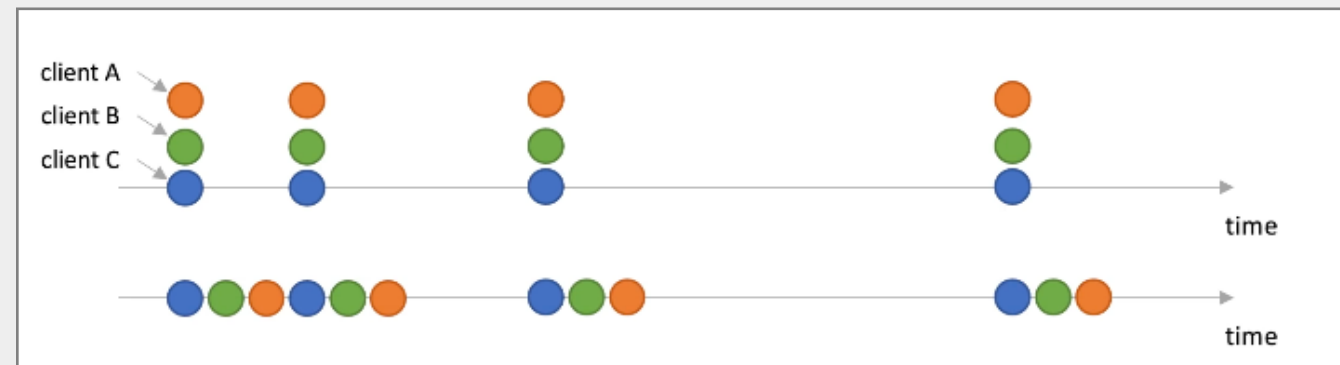


HTTP Status Code	Description	Retry?	Comments
N/A	Network connection is not established.	Yes	It doesn't matter if the API is idempotent or not. If no connection was established, the request was never sent.
400	Bad request.	No	There is a problem with the request. For example, a required parameter is not supplied for the specified action. Retry will not help.
403	Access denied.	No	One of many possible reasons - the provided security credentials are not valid.
404	Not found exception.	No	There is no requested resource present on the server side. E.g. the specified S3 bucket does not exist.
429	Too many requests.	Yes (but after a delay)	The request was not processed. It was immediately rejected by the server due to the high request rate.
500	Internal server error.	Yes (if idempotent)	The service encountered an unexpected condition that prevented it from fulfilling the request.
503	Service unavailable.	Yes	The requested service is not available.

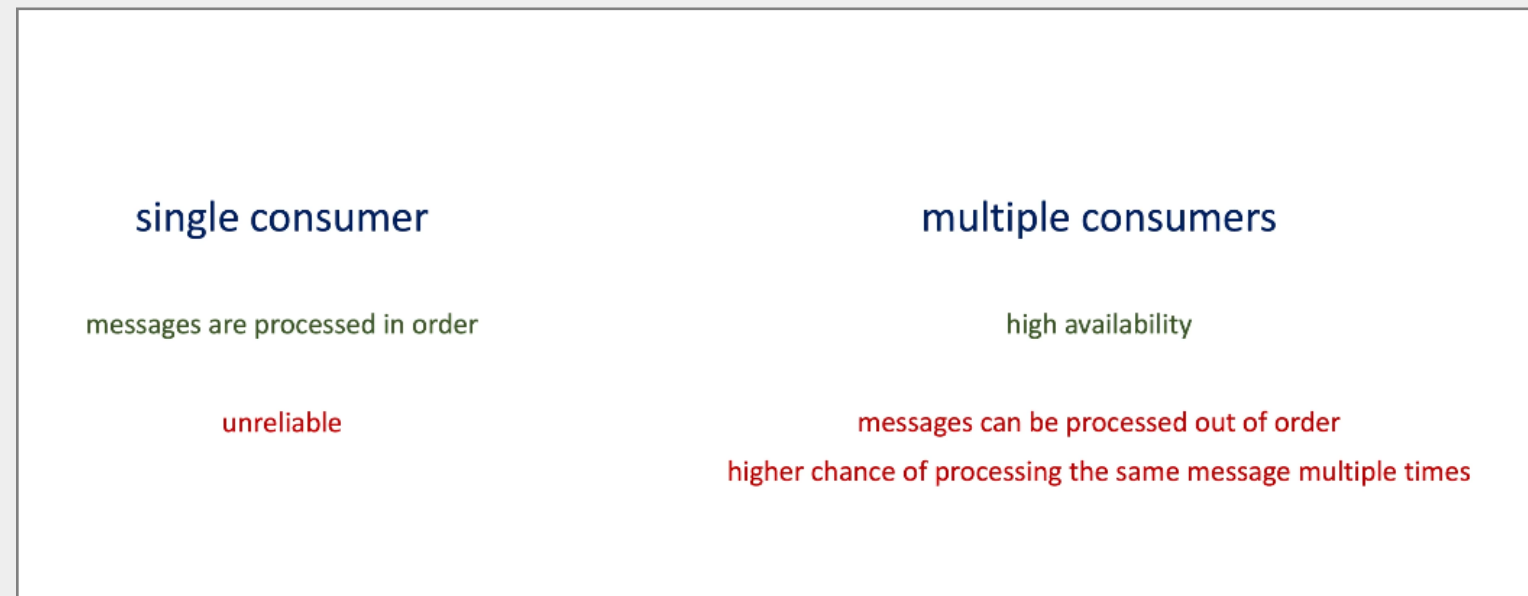
Exponential backoff



Retry at different times

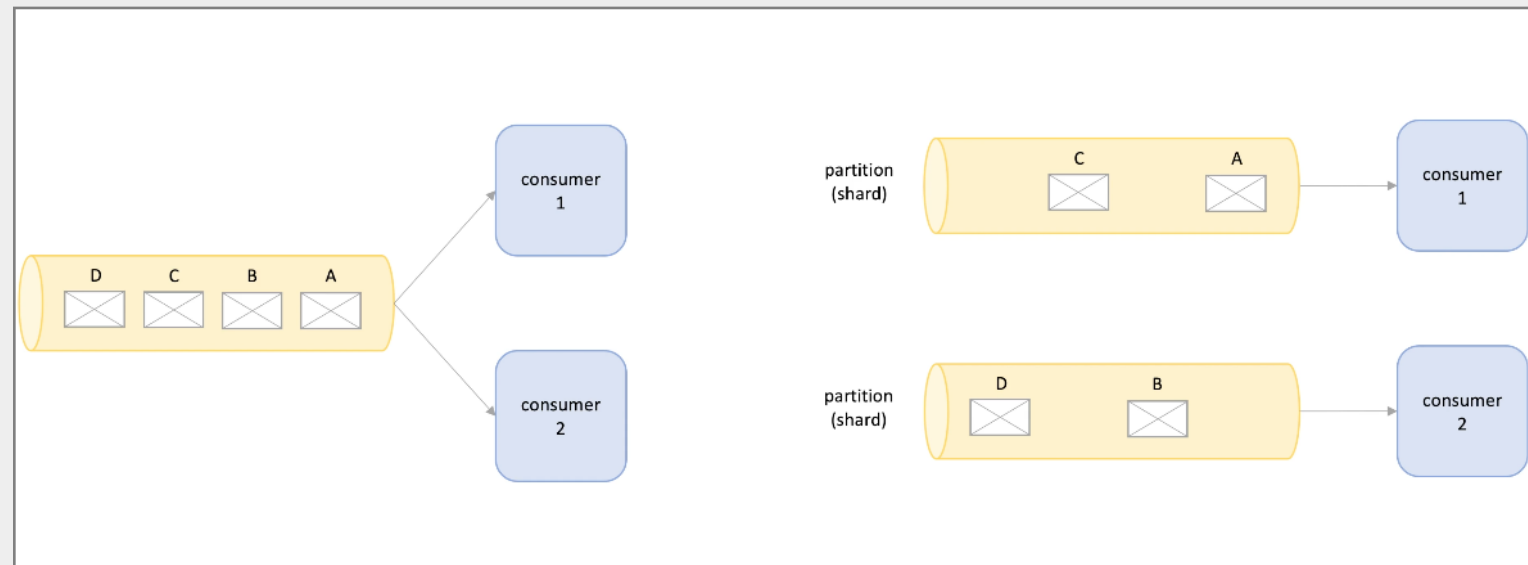


Single consumer vs Multiple consumers



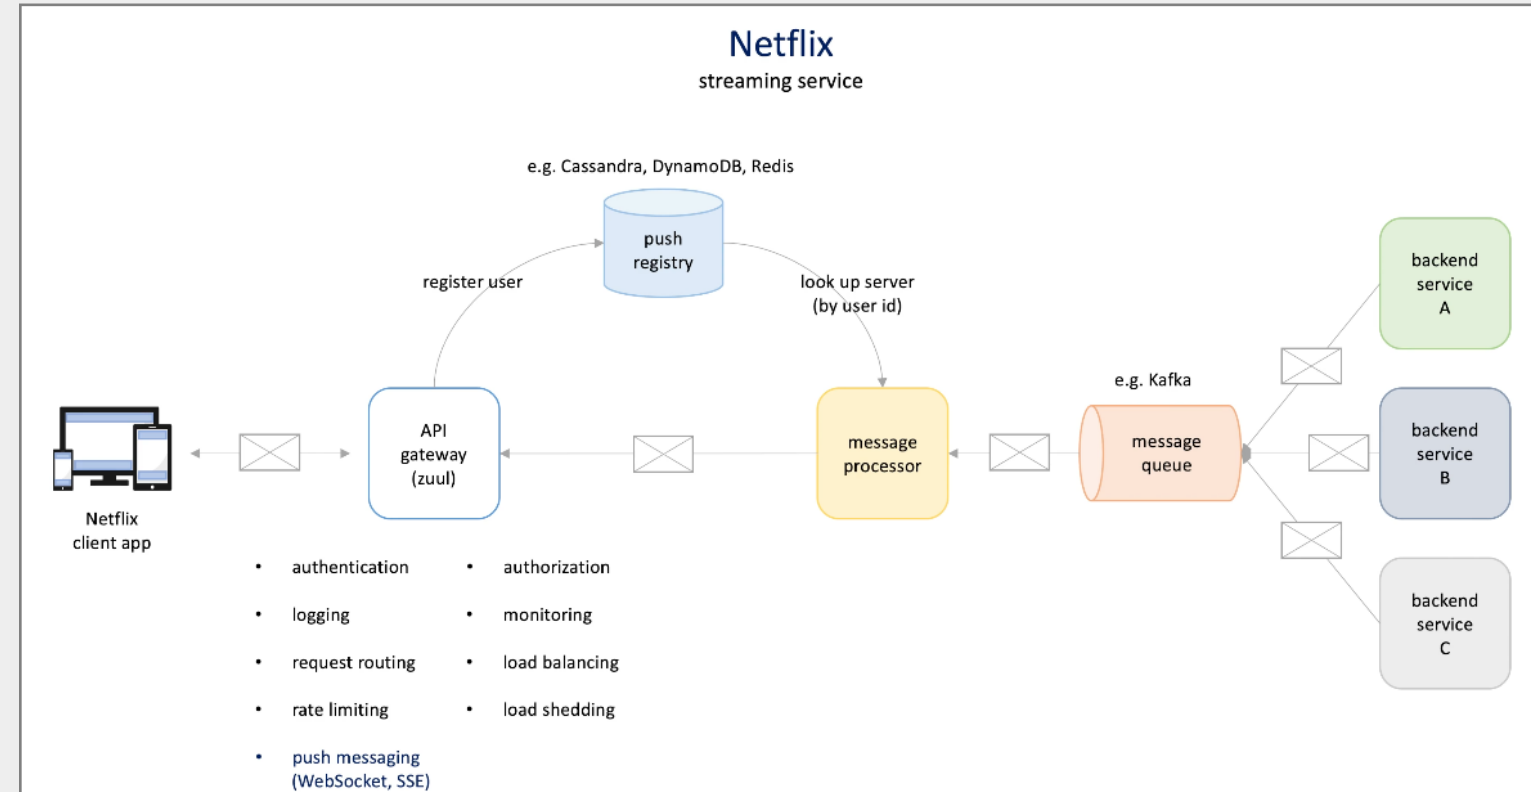
Multiple consumers break the order of messages

Data partitioning



[MQ Pub-Sub P&P Competing Claim Retry Partitioning]

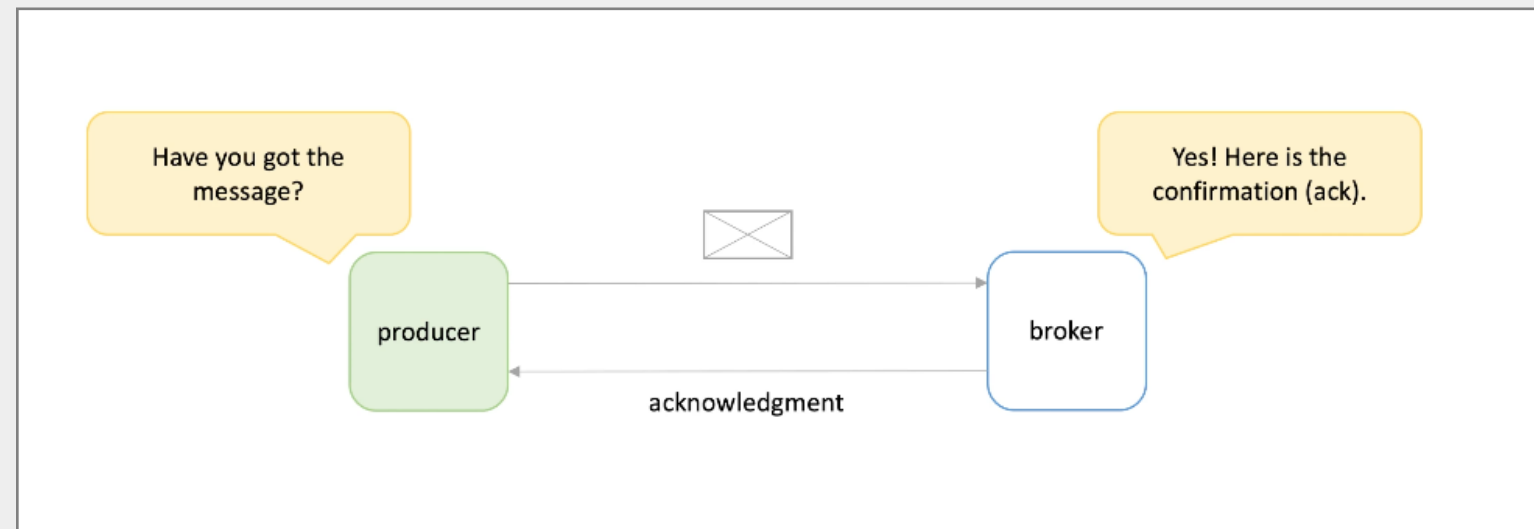
Netflix architecture



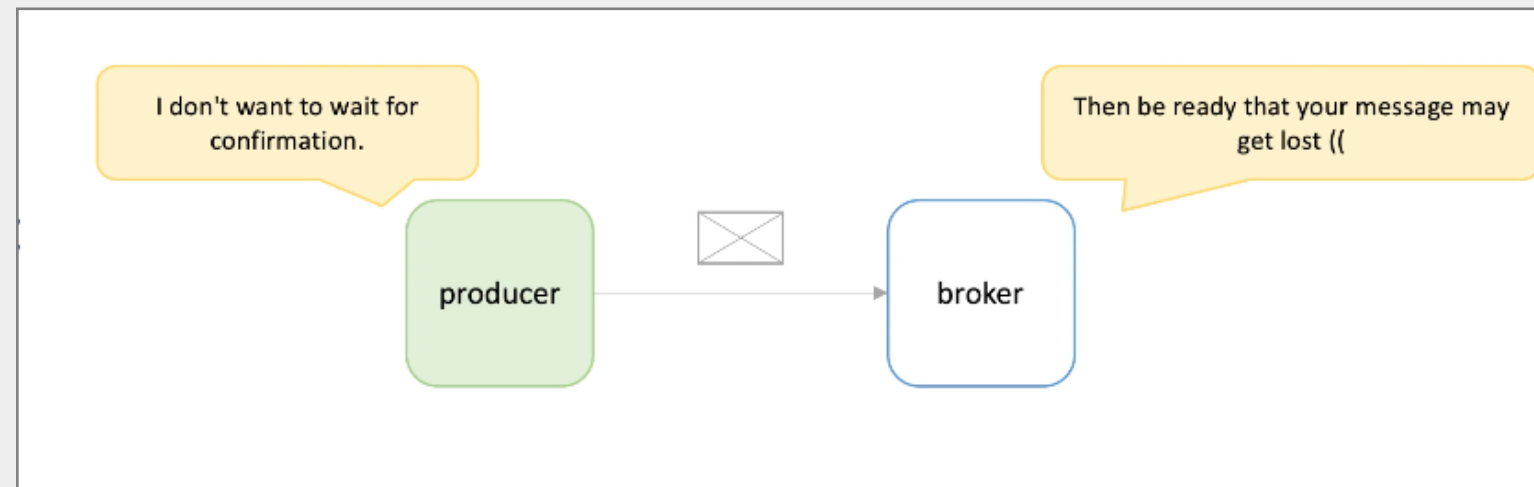
Chapter #3:

Internals of messaging

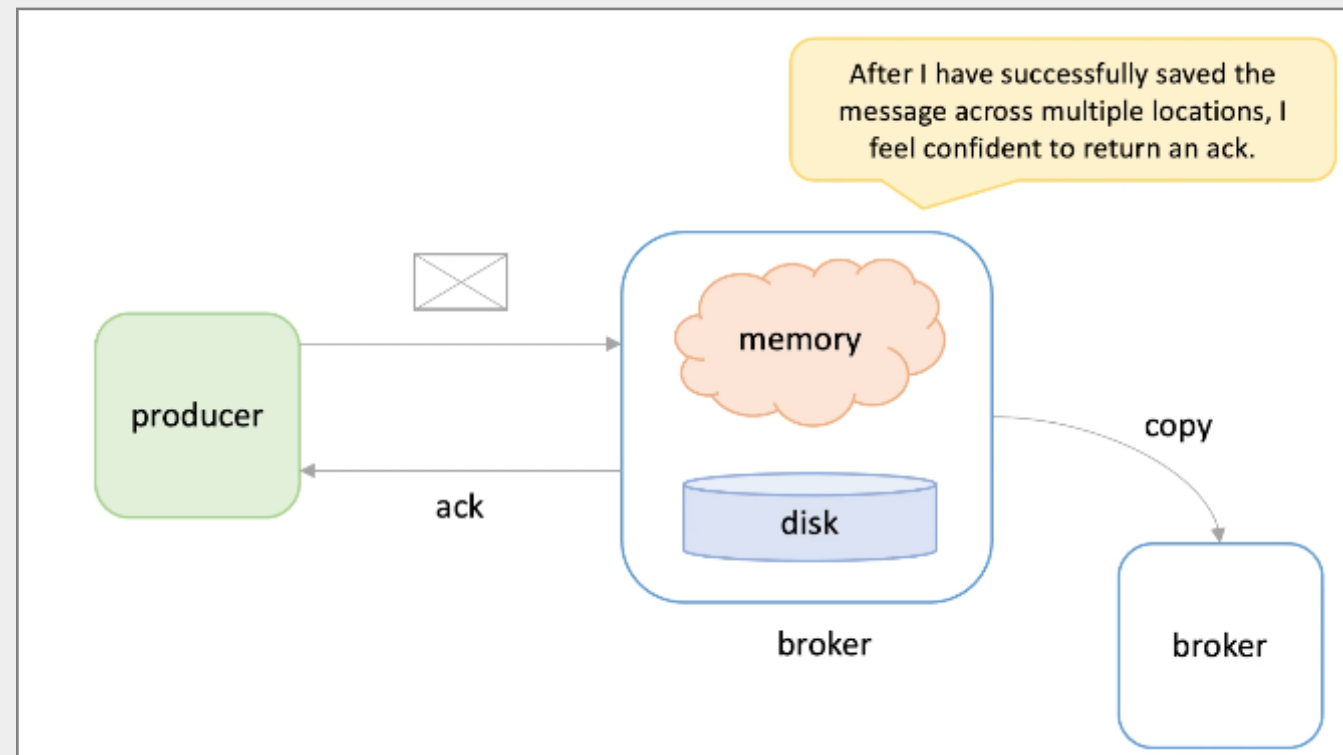
Message Acknowledgement



Unsafe ack aka “fire and forget”



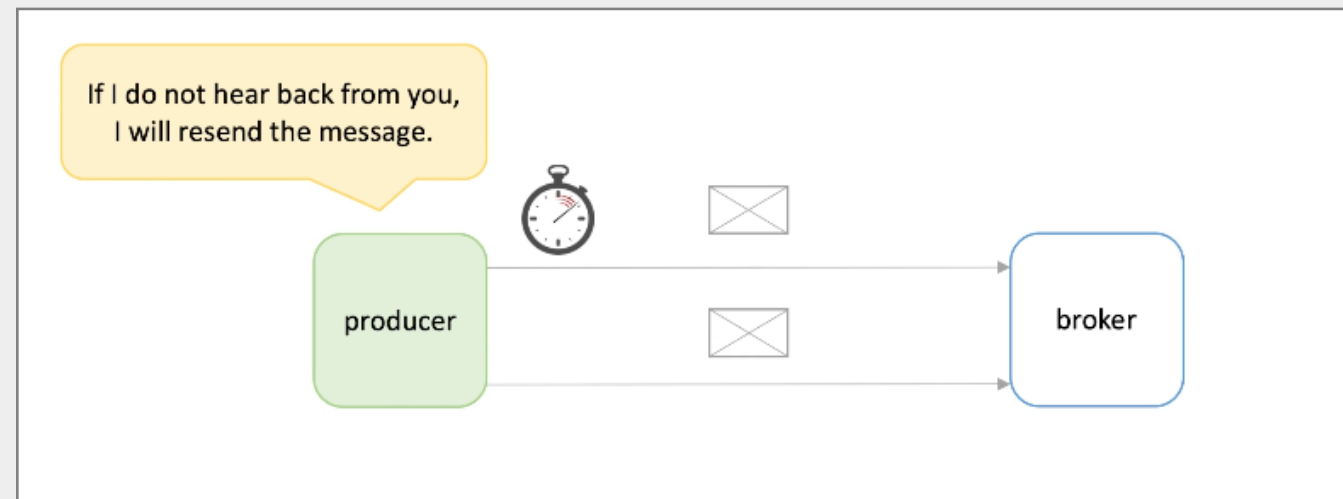
Safe ack



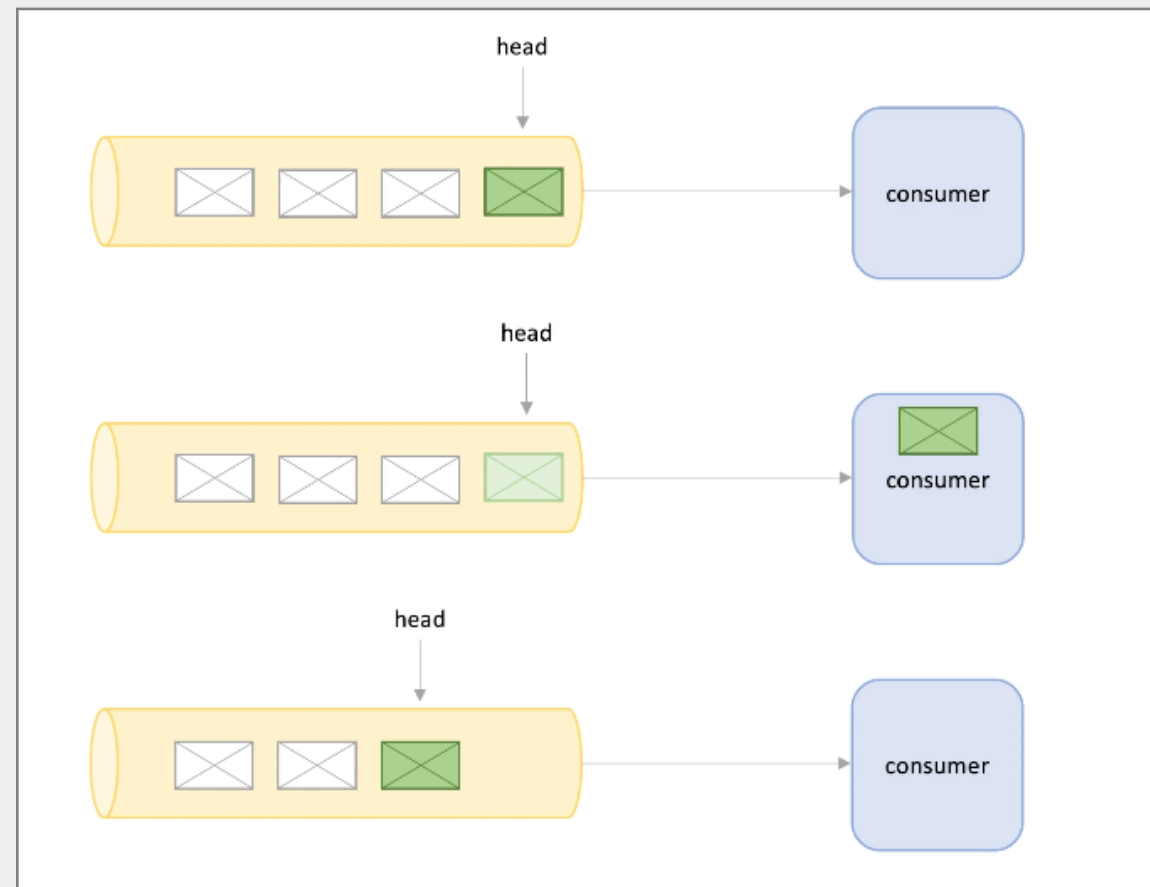
It is up to the broker what safe really means

[Ack Offset Polling Guarantees]

Failed ack



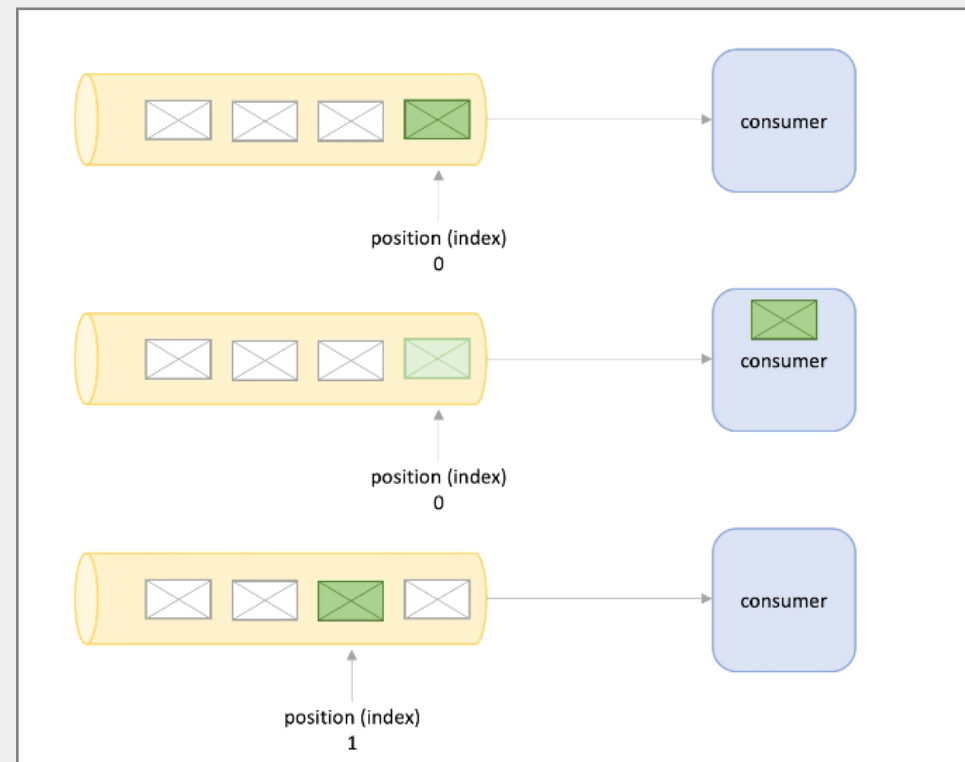
Consumer Offsets - Head delete



RabbitMQ, SQS

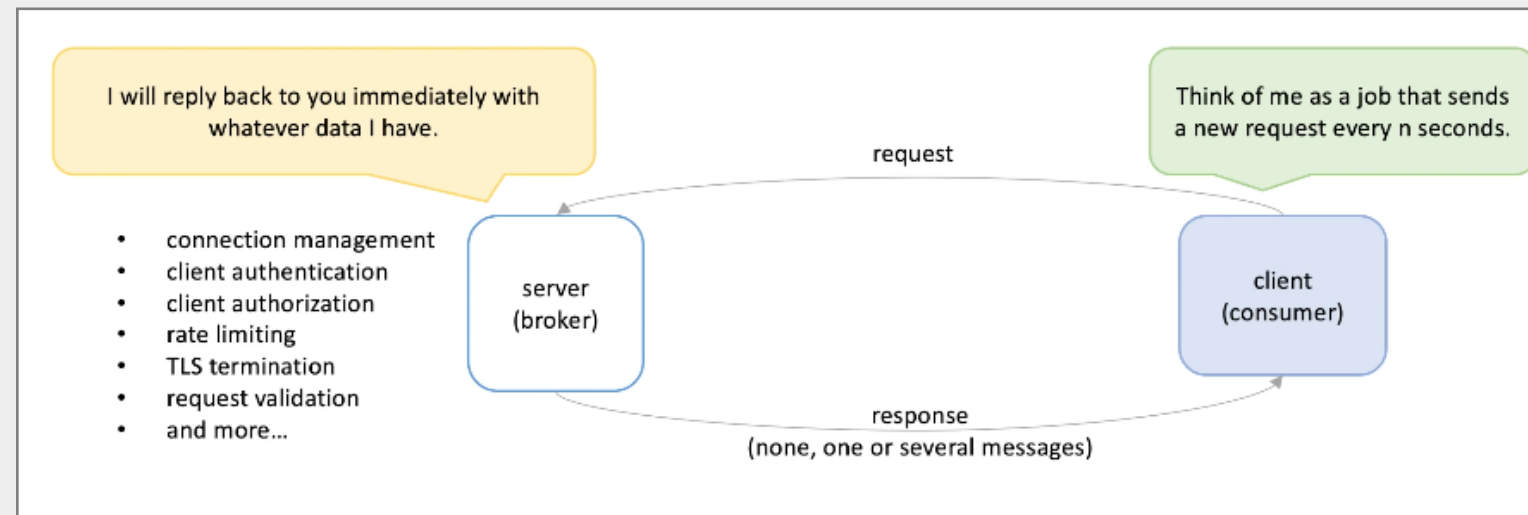
[Ack Offset Polling Guarantees]

Incrementing the pointer



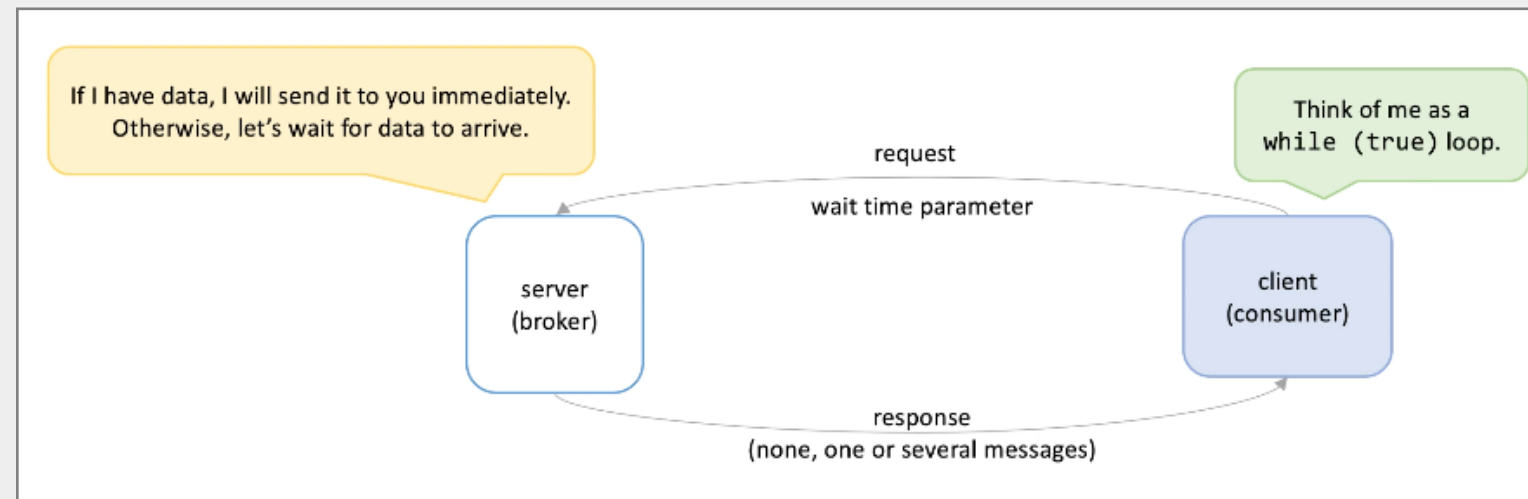
Kafka, Kinesis

Short Polling



[Ack Offset [Polling](#) Guarantees]

Long Polling



Delivery Guarantees

`at-most-once` -> Messages might be lost

`exactly-once` -> Messages processed only one time

`at-least-once` -> Messages might be processed multiple times

But failures will happen!



“ Failures are a given and everything will eventually fail over time: from routers to hard disks, from operating systems to memory units corrupting TCP packets, from transient errors to permanent failures. This is a given, whether you are using the highest quality hardware or lower cost components. ”

— Dr. Werner Vogels

Most of the messaging systems you will work with will be configured to provide `at-least-once` delivery guarantee